GENERAL ASSEMBLY

INTRO TO DATA SCIENCE

LECTURE 2: MACHINE LEARNING

LAST TIME:

- FIRST LOOK AT DATA SCIENCE & THE DATA MINING WORKFLOW
- DATA VISUALIZATION WITH R & GGPLOT2
- FIRST LINEAR MODEL

QUESTIONS?

What's big data?

The practical viewpoint:

1. $O(n^2)$ algorithm feasible: small data
2. Fits on one machine: medium data
3. Doesn't fit on one machine: big data

source: http://people.cs.umass.edu/~mcgregor/stocworkshop/langford.pdf

I. WHAT IS MACHINE LEARNING?
II. MACHINE LEARNING PROBLEMS

EXERCISES:
III. MULTIPLE REGRESSION & FEATURE EXTRACTION

I. WHAT IS MACHINE LEARNING?

from Wikipedia:

“Machine learning, a branch of artificial intelligence, is about the construction and study of systems that can *learn from data*.”

source: http://en.wikipedia.org/wiki/Machine_learning

WHAT IS MACHINE LEARNING?

from Wikipedia:

“Machine learning, a branch of artificial intelligence, is about the construction and study of systems that can *learn from data*.”

“The core of machine learning deals with *representation* and *generalization*…”

- *representation* – extracting structure from data
- *generalization* – making predictions from data

source: http://en.wikipedia.org/wiki/Machine_learning

II. MACHINE LEARNING PROBLEMS

supervised	making predictions
unsupervised	extracting structure

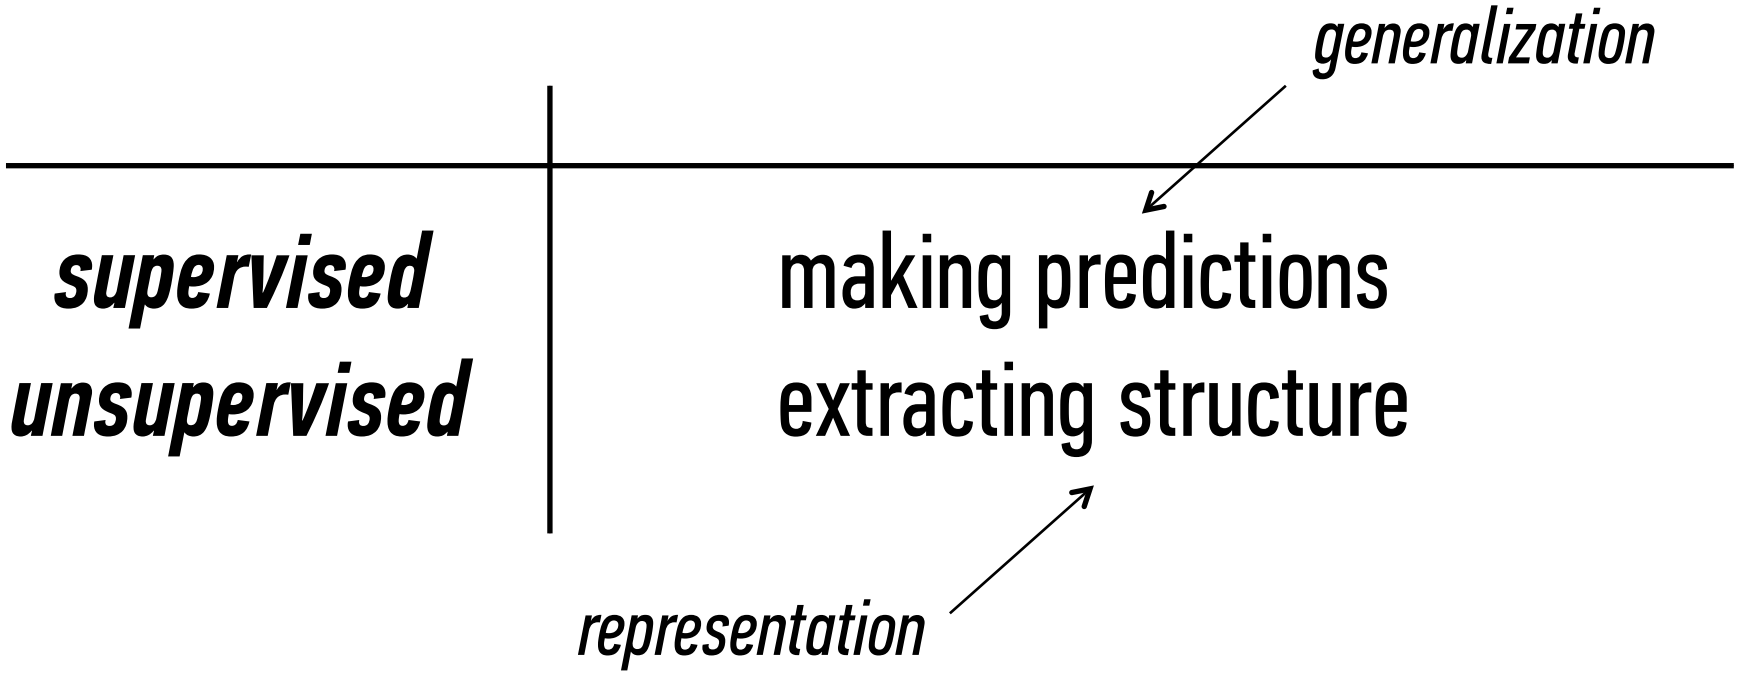
generalization
supervised
unsupervised
making predictions
extracting structure
representation

	continuous	*categorical*
	quantitative	qualitative

	continuous	*categorical*
	quantitative	qualitative

NOTE

The space where data live is called the *feature space*.

Each point in this space is called a *record*.

	continuous	*categorical*
supervised	regression	classification
unsupervised	dimension reduction	clustering

	continuous	*categorical*
supervised	regression	classification
unsupervised	dimension reduction	clustering

NOTE

We will implement solutions using *models* and *algorithms.*

Each will fall into one of these four buckets depending on the type of problem and type of data.

GENERAL ASSEMBLY

DATA SCIENCE AND MACHINE LEARNING

REMEMBER THIS?

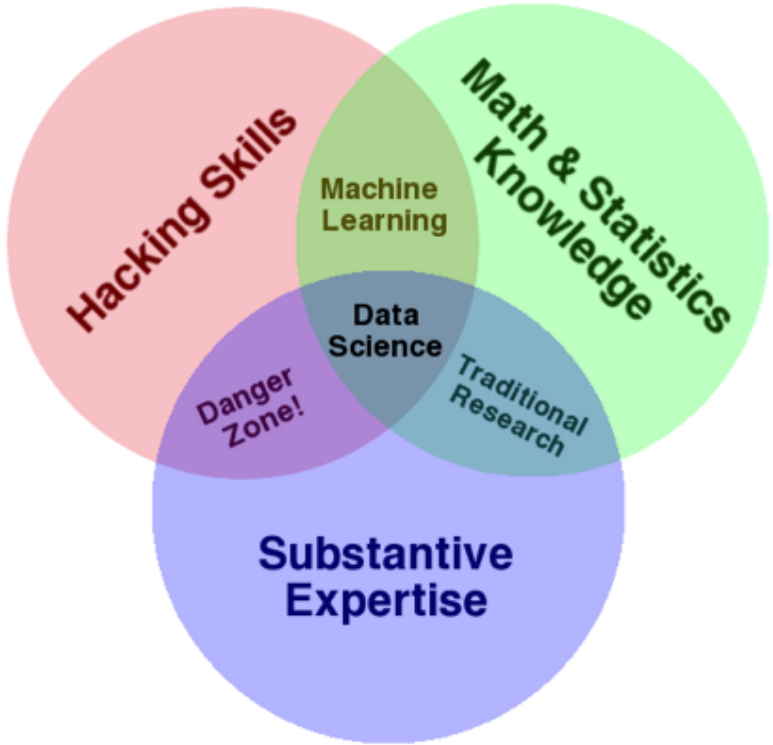

source: http://www.dataists.com/2010/09/the-data-science-venn-diagram/

YOU ARE HERE

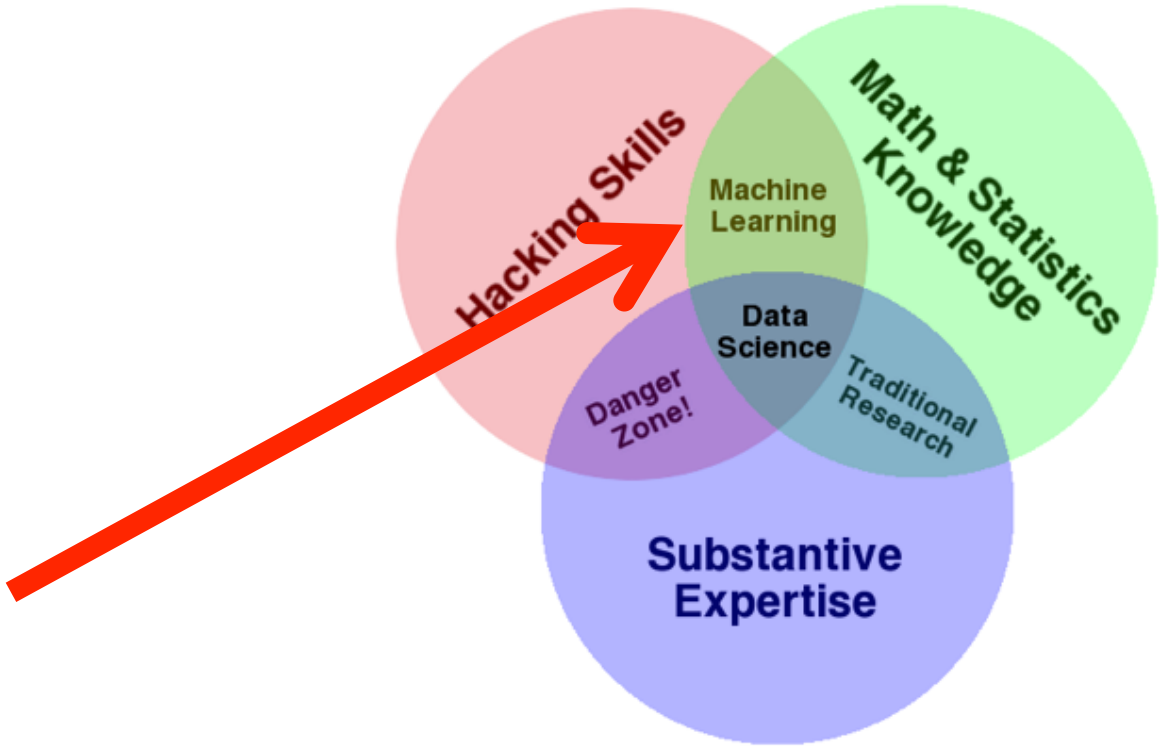

source: http://www.dataists.com/2010/09/the-data-science-venn-diagram/

YOU WANT TO GO HERE

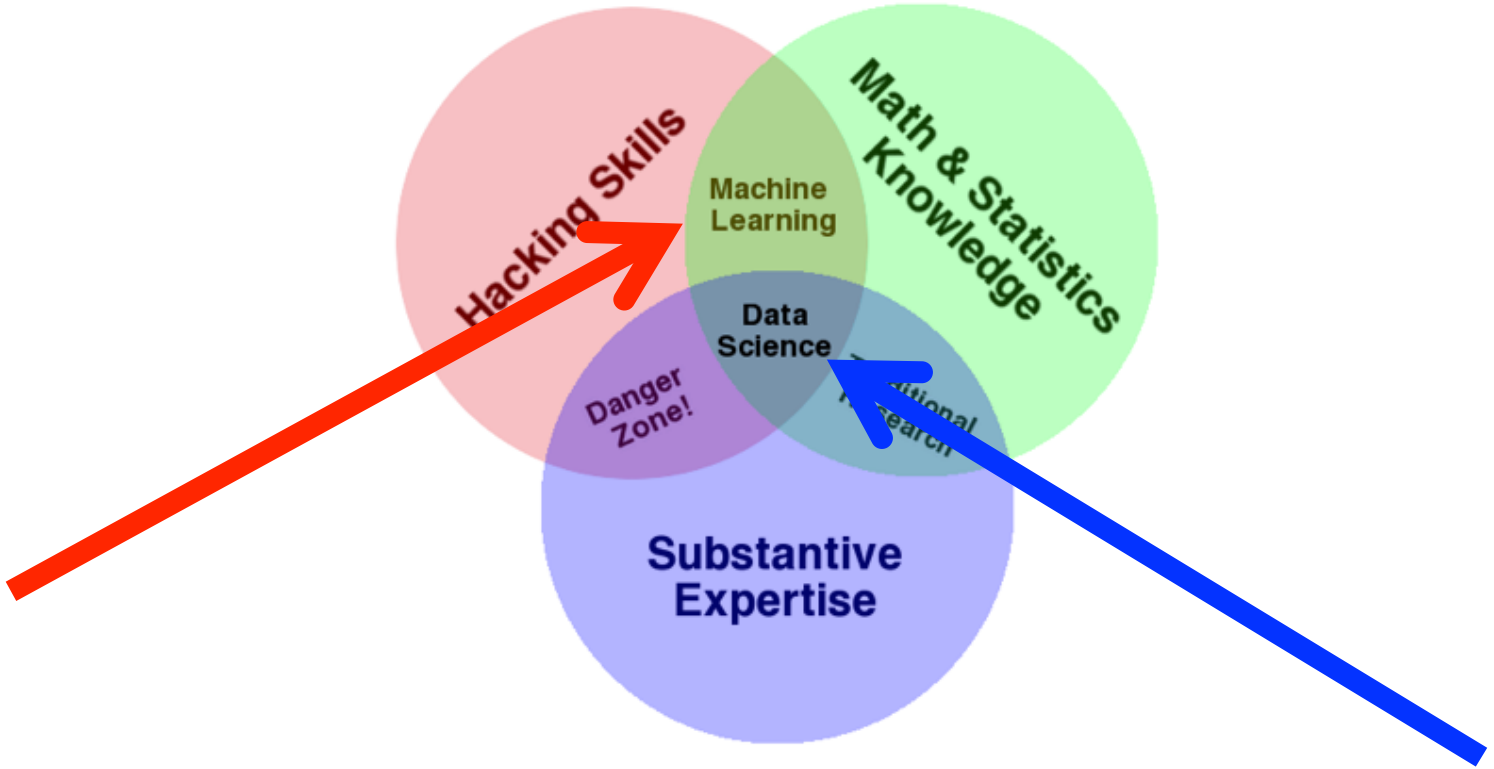

source: http://www.dataists.com/2010/09/the-data-science-venn-diagram/

YOU WANT TO GO HERE

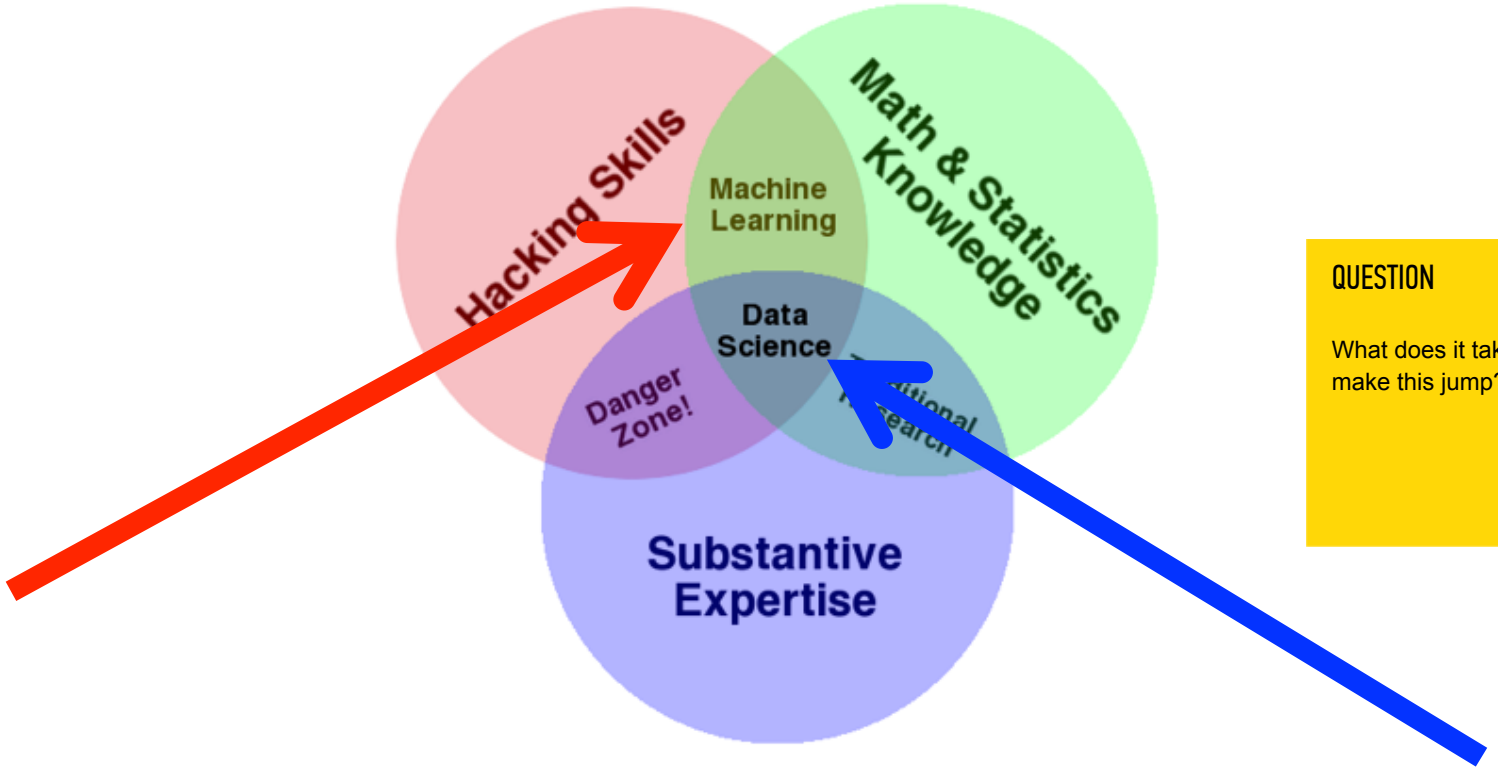

QUESTION

What does it take to make this jump?

source: http://www.dataists.com/2010/09/the-data-science-venn-diagram/

ANSWER: PROBLEM SOLVING!

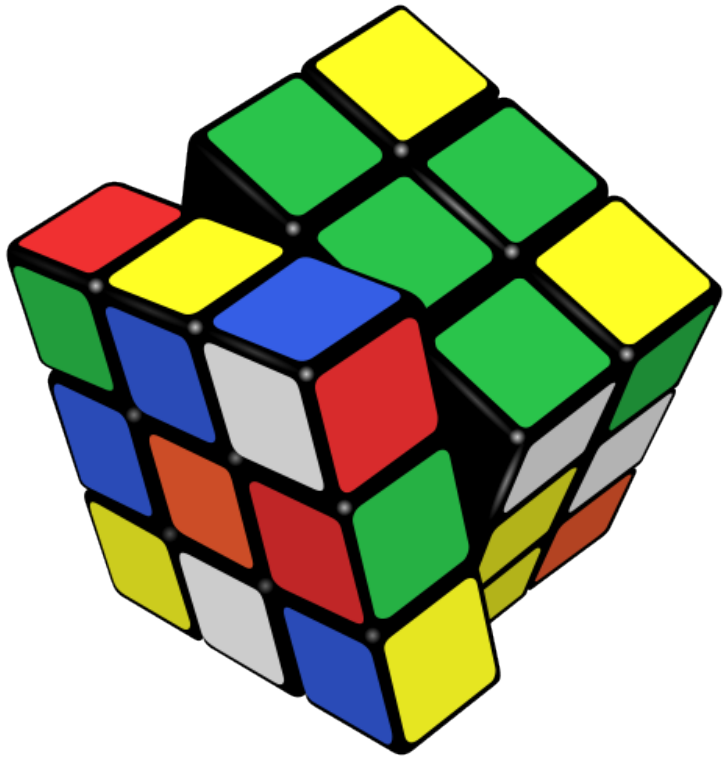

NOTE

Implementing solutions to ML problems is the focus of this course!

III. RELATIONSHIPS AMONG SEVERAL VARIABLES

EXERCISE – MULTIPLE REGRESSION (BACKWARD ELIMINATION)

KEY OBJECTIVES

- Create a regression model using several independent variables

- Extract meaningful features

TOOLS

- R (plot, lm, update)

1) **Linearity** of the relationship between dependent and independent variables (doesn't mean the relation between y and x has to be linear since we can use transformations if y and x as well)
2) **Independence** of the errors
3) **Homoscedasticity** (constant variance of the errors)
    1) versus time
    2) Versus the predictions or any independent variables
4) **Normality** of the error distribution

DISCUSSION